The Shadow Price of Neutrality

TODO(Isaiah): author

***Abstract*—TODO(Isaiah), Abstract, 0.2 page (claim → evidence):**

- **Problem: a long/short pod under a platform's neutrality mandate at a fixed 3%/yr risk budget → Sec. II.**
- **Approach: duals of the convex program at a root-found λ^*; three integer sensitivity estimators against certified secant intervals; big-M against perspective formulations → Sec. III.**
- **Q1 finding: the liquidity cap, sector bounds and trading costs carry the cost; factor neutrality is nearly free for α_{mom}; at zero, sector and all-style neutrality cost about the same → 0.16, 0.014, 0.16, 0.961 against 0.956, 0.964; Figs. 1–2.**
- **Q2 finding: the fixed-integer dual is exact when neither the support nor the active set changes within the step (max error 6.7×10^{-12}) and lies in $[L, H]$ in 4 of 23 binding cases; the best estimator depends on the step: at 2017-12 ($n = 11$) the fixed-integer dual is best at $\delta = 5\%$ (0.014; perspective 0.052, big-M 0.24) and the perspective slope at 10% (0.016 against 0.023; big-M 0.073) → Table III.**
- **Q3 finding: two mechanisms. The default-scaling root bound is consistent with Gurobi's outer approximation stopping early (29 of 36; 0.015 against 5.7×10^{-4} with a QCP node relaxation); the cone tolerance inflates incumbents (6.4×10^{-3}). The rescaled formulation (rows and cones rescaled together) removes both; presolve, not the formulation, closes big-M's gap → 3.3×10^{-4}, 0.015, 6.0×10^{-4} against 0.057; Fig. 4.**

I. Introduction

TODO(Isaiah), I, 0.6 page (claim → evidence):

- Motivation and significance: a multi-manager platform imposes neutrality to protect its other books; each constraint costs the pod expected alpha; the dual is the price of that constraint → Sec. II.
- Related work: fundamental law and transfer coefficient [1], [2]; constraint attribution [3]; factor alignment [4], [5]; realistic shorts [6]; perspective reformulations [7], [8]; integer sensitivity [9], [10]; sparse portfolios [11]; course texts [12], [13].
- Questions Q1–Q3 → Sec. II-G.
- Contribution: costs at equal risk via duality at λ^*, with the alignment closed form $\sqrt{1-\varphi}$ → A1–A4, Fig. 2.
- Contribution: certified secant intervals that test integer sensitivity estimators against solver error → A8, Table III, 3.6×10^{-4}.
- Contribution: a support-stability account of when the fixed-integer dual is exact → 6.7×10^{-12} (P1), 0.056 (P3).
- Contribution: the perspective model's two scaling mechanisms kept apart, and the rescaled formulation (rows and cones together) as the fix → A9, Sec. IV-C, 29 of 36, 1.9×10^{-11}.
- Project type: problem formulation, reusing published formulations in a new setting.

II. Problem Setting

A. Decision problem

A *pod*, one long/short equity team at a multi-manager platform, receives each month a forecast α_i of each stock's (*name*'s) *alpha*, its expected return beyond what its factor exposures explain, and chooses weights w as fractions of AUM (assets under management, \$1B). The platform caps the *ex-ante volatility*, the risk model's forecast $\sqrt{w^\top \Sigma w}$ annualized, at 3%/yr ($\sigma_m = 3\%/\sqrt{12}$ per month) and imposes a mandate that protects its other books. The *factor exposure* is $y = B^\top w$, B the loadings on the market and five styles (size SMB, value HML, profitability RMW, investment CMA, momentum Mom); *beta* is y_{mkt}. *Dollar*, *beta*, style and *sector neutrality* bound the net weight $|1^\top w|$, $|y_{\text{mkt}}|$, each $|y_k|$ and each $|S_j^\top w|$, S_j the 0/1 indicator of GICS sector j (the standard industry classification; 11 sectors). *Gross leverage* $1^\top(w^+ + w^-)$ is the long plus the short weight. A stock's *ADV* is its median daily dollar volume over 63 trading days; the *liquidity cap* is $u_i = \min(u, \theta\,\text{ADV}_i/\text{AUM})$. Shorts pay a *borrow cost* c per month (per year in Table I) and trades a one-way *trading cost* κ (1 bp $= 0.01\%$); from cash ($w_0 = 0$) trading costs κ times gross. Table I lists the base mandate.

B. Data and risk model

Daily adjusted prices and volumes of the current S&P 500 constituents (500 tickers, 2005-01-04 to 2025-12-31) come from yfinance, sectors from the constituent list, and the daily Fama–French five factors, momentum and risk-free rate from the French library [14]–[16]. A day is valid for a stock with a return and a dollar volume of at least \$1M; a stock is *eligible* at a date if at least 80% of the 252 trading days ending there are valid. At each of the 241 month-ends (2005-12 to 2025-12) eligible stocks' daily *excess returns* (minus the risk-free rate) over that window are regressed on the six factors: B holds the loadings, $D = \text{diag}(d)$ the residual variances and F the sample factor covariance, both scaled by 21 trading days to monthly, with d floored at its 5th cross-sectional percentile, so $\Sigma = BFB^\top + D \succ 0$. The eligible universe has 373–499 names.

C. Alphas and alignment

The momentum score, each stock's return over the 252 trading days before the date skipping the last 21 (it exists at 240 of the 241 month-ends), is z-scored across eligible stocks, winsorized at ± 3 and re-standardized; the alpha is $\alpha_i = \text{IC} \cdot \omega_i z_i$ with information coefficient IC $= 0.03$ and

residual volatility $\omega_i = \sqrt{d_i}$ [1], an input signal (no backtest). A D^{-1}-weighted regression of an alpha α on the loadings splits it into a residual $\alpha_\perp = \alpha - B(B^\top D^{-1}B)^{-1}B^\top D^{-1}\alpha$ and a factor-spanned part $\alpha_\parallel = \alpha - \alpha_\perp$. The *alignment share* $\varphi(\alpha) = \alpha_\parallel^\top \Sigma^{-1}\alpha_\parallel / \alpha^\top\Sigma^{-1}\alpha \in [0,1]$, on α's own, unrescaled split, is the factors' part of $\alpha^\top\Sigma^{-1}\alpha$, the squared *information ratio* (the best ratio of expected alpha to ex-ante risk); rescaling α leaves it unchanged. α_{mom}'s $\alpha_\perp$ and $\alpha_\parallel$, each rescaled to α_{mom}'s information ratio, also serve as signals. The *transfer coefficient* $\mathrm{TC} = \alpha^\top w/(\sqrt{\alpha^\top\Sigma^{-1}\alpha}\sqrt{w^\top\Sigma w})$ is the Σ-metric form of Clarke et al.'s [2].

D. The convex problem P-RB

$$\begin{aligned}
&\text{With } w = w^+ - w^-,\ w^\pm \geq 0,\\
&\max\ \alpha^\top w - c\,1^\top w^- - \kappa\,1^\top(w^+ + w^-)\\
&\quad\text{s.t. } y^\top F y + \textstyle\sum_i d_i w_i^2 \leq \sigma_m^2, \quad y = B^\top w,\\
&\qquad |1^\top w| \leq \varepsilon_{\mathrm{net}},\ |y_{\mathrm{mkt}}| \leq \varepsilon_\beta,\ |y_k| \leq b_k,\ |S_j^\top w| \leq h_j,\\
&\qquad 1^\top(w^+ + w^-) \leq G, \quad w_i^+, w_i^- \leq u_i.
\end{aligned} \tag{1}$$

Since $c, \kappa > 0$, at most one of $w_i^\pm$ is positive at an optimum. The optimal value A^* is net expected alpha per month. Row k has multiplier μ_k; a parameter p (a bound, u or θ) has the *family dual* $\mu_p = dA^*/dp = \sum_k \mu_k\,\partial r_k/\partial p$ over the rows whose right-hand side r_k moves with p (A5): θ through every θ-limited cap ($\partial u_i/\partial\theta = \mathrm{ADV}_i/\mathrm{AUM}$), u through every cap equal to u, a common sector or all-styles bound through all its rows. A *key* is a parameter whose dual is reported: a family, or one sector's or style's row. The *elasticity* $e_p = \mu_p p/A^*$ ($A^* > 0$) is the percent change in alpha per percent change in p; a constraint *binds* if $e_p \geq 10^{-4}$. The risk row's multiplier λ^* is the price of risk at which a mean-variance investor picks the same portfolio (A1). A fixed risk budget, not a fixed λ, compares costs at equal risk: relaxing a constraint cannot buy alpha by levering up.

E. The integer problem P-MI

P-MI adds a cardinality limit K and a minimum position $\ell = 0.25\%$, moves the risk row into the objective at λ^*, and raises the position cap to $u = 6\%$:

$$\begin{aligned}
&\max\ \alpha^\top w - \lambda^*\big(y^\top F y + \textstyle\sum_i d_i w_i^2\big)\\
&\qquad - c\,1^\top w^- - \kappa\,1^\top(w^+ + w^-)\\
&\quad\text{s.t. the mandate of (1) without the risk row,}\\
&\qquad \ell z_i^\pm \leq w_i^\pm \leq u_i z_i^\pm,\ z_i^+ + z_i^- \leq 1,\ \textstyle\sum_i (z_i^+ + z_i^-) \leq K,\\
&\qquad z^\pm \in \{0,1\}^N.
\end{aligned} \tag{2}$$

λ^* is P-RB's multiplier on the same universe at the base mandate with $u = 6\%$, fixed for that date and universe: a Lagrangian relaxation of the risk row, with zero gap in the convex case. *Big-M* keeps $\sum_i d_i w_i^2$ with $M = u_i$; the *perspective* formulation replaces it by $\sum_i d_i q_i$ with rotated cones $w_i^2 \leq q_i\zeta_i$, $q_i \geq 0$, $\zeta_i = z_i^+ + z_i^-$; the *rescaled perspective* divides every weight, row and bound by $\max_i u_i$ and writes each cone in its name's bound, $t_i = w_i/u_i$, $q_i' = q_i/u_i^2$, $t_i^2 \leq q_i'\zeta_i$, which leaves the feasible set and objective unchanged but changes what Gurobi's absolute tolerances act on (A9). These are the three *arms* of N1–N3. With λ^* fixed, P-MI's ex-ante volatility drifts (from the stored solutions, N3): median 2.87%/yr in I1 (2.42–3.02, $n = 144$, 3 above the budget), 2.91% in I2 (2.85–2.96, $n = 82$), 2.77% in N1 (2.22–3.31, $n = 36$ per arm, 4 above), so Q2's slopes are of the risk-penalized objective. The *continuous baseline*, P-MI with $\ell = 0$ and no cardinality row, is mean-variance at λ^*: with the risk row binding it has P-RB's optimizer and constraint duals at the base point (A1), but the value $A^* - \lambda^*\sigma_m^2$.

TABLE I
BASE MANDATE (P-RB) AND P-MI PARAMETERS; STYLE BOUNDS IN LOADING UNITS.

Parameter	Base
Risk budget (ex-ante volatility)	3%/yr
Gross leverage G	2
Net (dollar) $\varepsilon_{\mathrm{net}}$	$\pm$2%
Beta ε_β	$\pm$0.05
Each style b_k	$\pm$0.10
Each sector h_j (net)	$\pm$2%
Position cap u (P-MI)	3% (6%)
Liquidity θ	5% of ADV
AUM	\$1B
Borrow c	30 bp/yr
Trading κ (one-way)	5 bp
Minimum position ℓ (P-MI)	0.25%

F. Problem classes

P-RB is a convex quadratically constrained program (QCQP); its risk row is the second-order cone $\|(L_F^\top y, \sqrt{d}\circ w)\|_2 \leq \sigma_m$ with $F = L_F L_F^\top$. When $\lambda^* > 0$ its optimal w is unique: by A1, w also maximizes the mean-variance objective at $\lambda = \lambda^*$ over the mandate, which is strictly concave in w ($\Sigma \succ 0$, and the costs are convex and piecewise linear in w); the split $w^\pm$ is then unique because $c, \kappa > 0$. The big-M P-MI is a convex mixed-integer QP (every node relaxation a convex QP); the perspective P-MI is a mixed-integer second-order cone program.

G. Questions

Q1 (convex): at a fixed risk budget, which mandate constraints bind, what does each cost, and what determines that cost? **Q2** (integer): once a cardinality limit and a minimum position enter, how well do continuous sensitivity tools (relaxation duals, fixed-integer duals) predict the true marginal cost of a constraint? **Q3** (formulation and numerics): how do the big-M and perspective formulations compare as the mandate tightens, and how do solver tolerances interact with problem scaling?

III. Solution Approach and Analysis

A. Experiments

Table II defines the experiments: C1–C2 answer Q1, I1–I2 Q2, N1–N3 Q3. D12 is 12 month-ends (2008-10, 2009-03, 2011-08, 2013-06, 2015-08, 2017-12, 2018-12, 2020-03, 2020-11, 2022-06, 2023-03, 2025-04) spanning calm and

TABLE II
EXPERIMENTS; SOLVES: ROWS OF EACH RUN'S MAIN SOLVE TABLE.

exp.	instances	varied, arms	solves
C1	D12 (12 dates); all eligible names; α_{mom}, $\alpha_{\perp}$, $\alpha_{\parallel}$	base mandate, and the idealized mandates (i), (ii); envelope at $\delta = 1\%$ (0.1%)	108
C2	D12; all eligible; α_{mom} ($\alpha_{\perp}$, $\alpha_{\parallel}$ on factor families)	one family swept over its grid, the others at base	2040
I1	2009-03, 2017-12, 2020-03; $N = 250$; $f \in \{0.25, 0.5, 0.75\}$	4 bounds at $\delta = 10\%$ (5% at 2017-12); big-M and rescaled perspective	288
I2	2017-12 ($K = 67$); 2020-03 ($K = 79$); $N = 250$	sector bound at 41 points on $[0, 2b]$; as I1	164
N1	2017-12, 2020-03; $N \in \{100, 250, all\}$; $f \in \{0.25, 0.5\}$; 3 tightness levels	big-M, perspective (default scaling, rescaled); presolve off at $N = 250$; 3 reruns	147
N2(a)	N1 at $N = 250$, base tightness	the N1 arms at $s_{\mathrm{obj}} \in \{1, 10^4\}$	24
N2(b)	C1 base at 5 dates	Gurobi barrier, Clarabel, SCS; risk-row normalization and objective scaling on/off; reference	65
N3	the N1 instances at default presolve	root only, both perspective arms; MIQCPMethod -1, 0, 1 (default scaling also FeasibilityTol 10^{-9})	252

stressed markets. C1's 72 idealized solves are two mandates $\times$ three alphas $\times$ 12 dates: (i) all six exposures zero, (ii) the same plus dollar neutrality, each with the risk budget only. C2 sweeps one family at a time, the others at base, over (off first): dollar 0.1, 0.05, 0.02, 0.01, 0; beta 0.2, 0.1, 0.05, 0.02, 0; each style and all styles 0.4, 0.2, 0.1, 0.05, 0.02, 0; sector 0.08, 0.04, 0.02, 0.01, 0; gross 4, 3, 2, 1.5, 1; u 0.06, 0.03, 0.015; θ 0.1, 0.05, 0.025. N counts names (top N by ADV); $K = \mathrm{round}(fS)$, half up, S the support size of P-RB at $u = 6\%$. The design pilot's rule set the I1 bounds: the families with the largest median C1 elasticity at the I1 dates (u and θ excluded: they set the big-M constants) and the Mom bound at 0.025, kept if certified binding (A8) at $f = 0.5$ on 2017-12, at most 4.

B. Algorithms

P-RB is solved as the mean-variance QP $\max \alpha^\top w - \lambda w^\top \Sigma w -$ costs by Gurobi's dual simplex (tolerances 10^{-9}: a basic solution, its active set and duals exact up to them), with λ root-found by Brent's method on $\log \lambda$ until volatility meets the budget to relative 10^{-10} (A1). Clarabel's primal–dual interior-point method [17] (tolerances 10^{-12}) solves the cone form of (1) as a cross-check and the perspective relaxation of (2) ($z \in [0, 1]$), whose finite differences give its slope; dual simplex solves the big-M relaxation and the fixed-integer QP (P-MI with z fixed at the incumbent). *Plain relaxations* are the $z \in [0, 1]$ relaxations solved outside the MIP, without its presolve or root cuts. Gurobi's branch and bound [18] stops when the relative gap is below *MIPGap* (10^{-6} in I1, I2; 10^{-4} in N1, N2(a)) or its *Work*, a deterministic effort count, reaches $W^* = 340$ ($2W^*$ in I1, I2), the Work of 300 s at the pilot's median 1.12 units/s; MIQCPMethod (cone handling at the nodes) stays at its default. The objective is multiplied by $s_{\mathrm{obj}} = 10^4$, the power of ten that brings the pilot's proxies for $|U^*|$ (the mean-variance value at λ^* on each I1 and N1 universe) to 9.1–21; values are reported unscaled. The incumbent is *polished*: the fixed-integer QP at its z, resolved at tolerances 10^{-9}, gives V. N2(b) runs Gurobi's barrier and CVXPY's [19] Clarabel and SCS (ECOS is not installed) at default tolerances against Clarabel at 10^{-10}.

C. Analysis

A1 (risk budget and mean-variance). With multipliers λ (risk row) and ν (mandate), the Lagrangian of (1) is $\alpha^\top w -$ costs $- \lambda(w^\top \Sigma w - \sigma_m^2) - \nu^\top$(mandate); maximizing it over the mandate is mean-variance at λ. For $\lambda > 0$ its maximizer $w(\lambda)$ is unique ($\Sigma \succ 0$), with volatility continuous and nonincreasing in λ, so if the volatility at $\lambda_{\min} = 10^{-6}$ exceeds the budget a root λ^* exists. There the risk row is tight (primal feasibility, complementary slackness), and Slater's condition holds ($w = 0$ strictly satisfies the risk row; the other rows are affine), so KKT is sufficient: $w(\lambda^*)$ solves P-RB with multiplier λ^*. Conversely, by strong duality every P-RB optimum maximizes the Lagrangian at λ^* (used in Sec. II-F). Otherwise the budget is slack ($\lambda^* = 0$) and the solver returns $w(\lambda_{\min})$, feasible and with net alpha within $\lambda_{\min}\sigma_m^2 = 7.5 \times 10^{-11}$ of P-RB's optimum (4 of 36 C1 base solves, all $\alpha_{\parallel}$).

A2–A3 (equality case). With only the rows $Aw = 0$ (A of full row rank), no costs and no other rows, stationarity gives $w^* = \Sigma^{-1} g/(2\lambda)$ with $g = \alpha - A^\top \nu$; $Aw^* = 0$ gives $\nu = (A\Sigma^{-1}A^\top)^{-1} A\Sigma^{-1}\alpha$. If $g \neq 0$ ($\alpha \notin \mathrm{range}(A^\top)$) the risk row binds, so $\lambda = \sqrt{g^\top \Sigma^{-1} g}/(2\sigma_m)$ and $A^*_{\mathrm{eq}} = \alpha^\top w^* = \sigma_m \sqrt{g^\top \Sigma^{-1} g}$. Divided by the unconstrained $\sigma_m \sqrt{\alpha^\top \Sigma^{-1} \alpha}$, the retained fraction $\sqrt{g^\top \Sigma^{-1} g / \alpha^\top \Sigma^{-1} \alpha}$ equals TC at w^*, and it is at most one by Cauchy–Schwarz in the Σ^{-1} metric, since g is the Σ^{-1}-orthogonal projection of α onto $\{v : A\Sigma^{-1} v = 0\}$ along $\mathrm{range}(A^\top)$.

A4 (alignment). Let B have full column rank. $B^\top D^{-1} \alpha_{\perp} = 0$ (the regression residual) gives $\Sigma D^{-1} \alpha_{\perp} = (D + BFB^\top) D^{-1} \alpha_{\perp} = \alpha_{\perp}$, so $\Sigma^{-1} \alpha_{\perp} = D^{-1} \alpha_{\perp}$ and $B^\top \Sigma^{-1} \alpha_{\perp} = 0$, with no assumption on F: under the risk budget alone $\alpha_{\perp}$'s optimum, proportional to $\Sigma^{-1}\alpha_{\perp}$, has zero factor exposure, and factor bounds $b_k > 0$ have zero shadow price. With $A = B^\top$, $g = \alpha_{\perp}$: the projection of A2 is unique, and $\alpha_{\perp}$ satisfies both $B^\top \Sigma^{-1} \alpha_{\perp} = 0$ and $\alpha - \alpha_{\perp} = \alpha_{\parallel} = B\gamma \in \mathrm{range}(B)$ (γ the regression coefficients). The cross term $\alpha_{\perp}^\top \Sigma^{-1} \alpha_{\parallel}$ vanishes, so $\alpha^\top \Sigma^{-1} \alpha = \alpha_{\perp}^\top D^{-1} \alpha_{\perp} + \alpha_{\parallel}^\top \Sigma^{-1} \alpha_{\parallel}$, and full factor neutrality retains $\sqrt{1 - \varphi(\alpha)}$ of the alpha.

A5 (envelope). By Slater's condition (A1) strong duality holds. A^* as a function of the right-hand sides is the perturbation function of a convex program, hence concave, and nondecreasing in each bound because loosening it enlarges the feasible set [12]. Where the multipliers are unique (for example under LICQ), $dA^*/dr_k = \mu_k$, also across nondegenerate active-set changes, where μ_k has a kink but A^* does

not; elsewhere only one-sided slopes exist, as at a cap tie $u = \theta\,\mathrm{ADV}_i/\mathrm{AUM}$. The chain rule gives the family dual of Sec. II-D.

A6 (perspective tightness). Over the same (w, z), restricted by the same linking rows, $z \in (0, 1]$ gives $q \geq w^2/z \geq w^2$ (and $q = w = 0$ at $z = 0$), so with $\lambda^* \geq 0$ the perspective relaxation's risk term is at least big-M's and its value at most big-M's.

A7 (one-asset toy). For $\max\, aw - \lambda d w^2$ ($\lambda d > 0$) over $w \in \{0\} \cup [\ell, u]$, $w \leq b$, the value is 0 for $b < \ell$ and jumps at $b = \ell$ by $\max(0, a\ell - \lambda d \ell^2)$, positive iff $a > \lambda d \ell$. Below ℓ the true local slope is 0, and so is the fixed-integer dual. The big-M relaxation's slope is $a - 2\lambda d b$ for $b < w^\circ = a/(2\lambda d)$ and 0 otherwise; the perspective relaxation's (q minimized at $z = w/\ell$, so $q = \ell w$) is $a - \lambda d \ell$ if $a > \lambda d \ell$ and 0 otherwise. For $\ell/2 < b < \ell$, $a - 2\lambda d b < a - \lambda d \ell$: where both are positive, big-M's slope is the closer to 0 although the perspective bound is the tighter. A tighter bound need not give a better slope.

A8 (integer value functions). The P-MI value $U^*(b)$ is the upper envelope, over support patterns, of concave pieces: nondecreasing by feasible-set inclusion, in general neither concave nor continuous. The relaxation dual is the slope of an upper-bounding concave function, the fixed-integer dual the slope of the active piece, and a re-solve gives a secant across pieces. With V the polished value, $\bar{U}$ Gurobi's final bound and $\bar{U}_{\text{eff}} = \max(\bar{U}, V) + 10^{-6}|V|$, $V \leq U^* \leq \bar{U}_{\text{eff}}$ at $b \pm \delta$ (δ absolute here, a percentage of b in the experiments) gives the certified interval

$$L = \frac{V_+ - \bar{U}_{\text{eff},-}}{2\delta} \leq \frac{U^*_+ - U^*_-}{2\delta} \leq \frac{\bar{U}_{\text{eff},+} - V_-}{2\delta} = H, \quad (3)$$

with point slope $s = (V_+ - V_-)/(2\delta)$. "Certified" is conditional on V's weights being feasible (the independent checker passes them at 10^{-9}) and on Gurobi's bound being valid up to the slack, a margin, not a proof. I1 has no violation (0 at 144 points); N2(a) at $s_{\text{obj}} = 1$ has a counter-case ($9.1 \times 10^{-5}|V|$, Sec. IV-C). A case is *certified binding* iff $L > 0$ (unlike *binds*, $e_p \geq 10^{-4}$, in the convex problems); $H < 0$ flags solver error. In I1 and I2 both formulations are solved at every point; V is the larger polished value, $\bar{U}$ the smaller bound, and the *certified gap* is $(\bar{U}_{\text{eff}} - V)/|V|$.

A9 (scaling). An absolute tolerance τ on a quantity of size m is a relative error τ/m; rescaling a row or the objective changes the effective tolerances of an equivalent problem. At an optimum with $\lambda^* > 0$ a held name has $q_i = w_i^2 \in [\ell^2, u_i^2]$ (q_i is minimized, not bounded), so Gurobi's absolute feasibility tolerance 10^{-6} is a relative error of up to $10^{-6}/\ell^2 = 0.16$ on the smallest cones, and at most $10^{-6}(u_i/\ell)^2 = 5.8 \times 10^{-4}$ with each cone in its name's bound. These scale a cone residual; they do not bound the objective error. They account for the incumbents' cone violation, with objective effect $\lambda^* \sum_i d_i \max(0, w_i^2/\zeta_i - q_i)$ reported relative to $|V|$, and for certified gaps above MIPGap: the default-scaling N2(a) solves at $s_{\text{obj}} = 10^4$ ended OPTIMAL (4 of 4), not at the WorkLimit, so their gap comes from the cone effect (Sec. IV-C). They do not account for the default-scaling root bound, which N3 traces to Gurobi's outer approximation stopping before it converges.

D. Finite differences and noise

The perspective relaxation's slope at b is the central difference $\hat{s} = (R(b + \delta_r) - R(b - \delta_r))/(2\delta_r)$, $\delta_r = 10^{-3}\,b$, of its values R: the objective at Clarabel's (w, ζ) with each q_i set to w_i^2/ζ_i (0 where $w_i = 0$), so every cone holds exactly. A relaxation fails if its residuals exceed 10^{-8} or its duality gap exceeds $\varepsilon = 10^{-9}\max(|R|, |R_{\text{ref}}|)$, R_{ref} the value at the unperturbed bound. If every stencil value is within ε of exact, noise moves $\hat{s}$ by at most $\eta = \varepsilon/\delta_r$; the gap test does not prove this premise, so η is a noise scale. The one-sided slopes $s_+ = (-3R_0 + 4R_1 - R_2)/(2\delta_r)$ and $s_- = (3R_0 - 4R_{-1} + R_{-2})/(2\delta_r)$, $R_j = R(b + j\delta_r)$, are recorded; $s_+ - s_-$ has absolute coefficients summing to 16 over $2\delta_r$, so under the same premise noise moves it by at most 8η, and a case is flagged as a kink if $|s_+ - s_-| > \max(10^{-4}|\hat{s}|, 8\eta)$. A relaxation's *gap* at b is $(R - V)/|V|$.

E. Validation

530 passing tests and logged checks with negative controls confirm the analysis: the QP and the Clarabel QCQP agree on A^* to 1.2×10^{-10} and on λ^* to 3.8×10^{-8} (A1); A2–A3 hold to 8.9×10^{-16}, the A4 split to 2.1×10^{-16}; on 202 one-asset toys (K29c's 200 random ones among them) both formulations match A7 at 8279 grid points to 1.4×10^{-17} (3 points within Gurobi's tolerance below ℓ excluded); P-MI with $K \geq N$, $\ell = 0$ reproduces P-RB's weights and A^* to 1.5×10^{-15} (largest relative error, in the weights). An independent feasibility checker passes every C1 and C2 solution and every polished V. The slope method was *gated* (a pass/fail check fixed in advance) before I1, on the random toys (0 misses at 8399 points: each grid point $b > 0$, ℓ and w°; 1 failed relaxation excluded) and 39 objective-scale pairs (0 misses). At $10 s_{\text{obj}}$ the I1 perspective slopes meet CHECKS K29b's tolerance at 48 of 48 cases (the 12 at 2009-03 as an acceptance check, the rest afterwards), moving at most 3.2×10^{-7} relative where certified binding ($n = 34$); none of the 240 I1 relaxations failed.

IV. Results

A. Q1: Constraint Costs in the Convex Problem

Setup. C1 and C2 (Table II) at the D12 dates on all eligible names; the criterion is the elasticity e_p from exact duals, validated as derivatives below; statistics are over D12 for α_{mom} ($n = 12$) unless stated.

Which constraints bind. At base, A^* has median 30.5 bp/month (range 22.0–38.3), a median 0.65 (0.51–0.77) of the unconstrained, costless optimum. The risk budget binds in all but 4 of the 36 base solves, all of them $\alpha_{\|}$. The liquidity cap binds at all 12 dates (median elasticity 0.16), sector bounds at 12 (0.014), dollar neutrality at 10 (1.3×10^{-3}), style bounds at 4 and the position cap at 2; beta and gross bind at none. Fig. 1 shows the three families whose curves move: at the tightest grid points the pod keeps a median 0.744 of its alpha

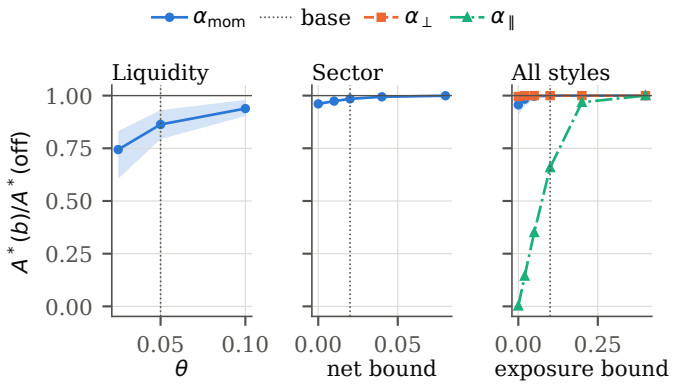


Fig. 1. Cost curves (C2), others at base (dotted: base); median and interquartile band over D12.

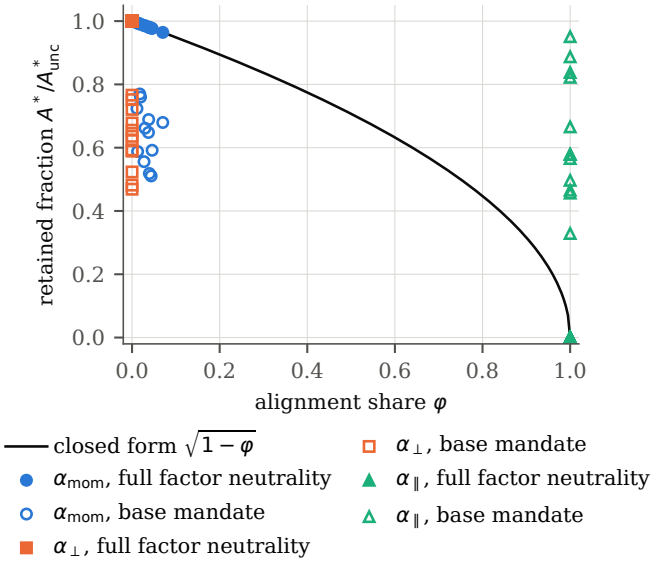


Fig. 2. Retained fraction against φ (C1), relative to the unconstrained, costless P-RB.

at $\theta = 2.5\%$ (range 0.52–0.95), 0.961 under exact sector neutrality and 0.956 with every style exposure at zero; every other family, the Mom bound among them, keeps at least 0.994 (gross $G = 1$).

What determines the cost. Factor bounds are cheap because α_{mom} is weakly aligned with the factors: φ is 0.011–0.070 over D12, and under idealized full factor neutrality the retained fraction is $\sqrt{1-\varphi}$ (A4), at least 0.964 (Fig. 2; numerical and closed-form A^* agree to 1.6×10^{-14} relative, the fractions to 7.5×10^{-8}, the square root of φ's rounding near $\varphi = 1$ for $\alpha_\parallel$, and to 1.9×10^{-15} otherwise). $\alpha_\perp$ ($\varphi \le 3.1 \times 10^{-33}$) loses nothing and $\alpha_\parallel$ ($\varphi = 1.0$) everything. At the base mandate the α_{mom} points lie well below the curve: the cost comes from the caps and the costs. Of the base portfolio's gross expected alpha $\alpha^\top w^*$, the pod keeps a median 0.80 (0.74–0.84) as A^*, pays 0.16 (0.13–0.21) in trading and 0.039 (0.032–0.052) in borrow.

Duals as derivatives. At every binding key and D12 date (115 checks over 18 keys: the binding families, every sector, and the CMA and Mom rows) the dual matches the central difference within 1% (largest error 3.1×10^{-4}) at $\delta = 1\%$ of the bound, or at 0.1% at the 64 flagged stencils (the dual zero at an end, or forward and backward differences more than 1% of the central one apart; 35 change the active set, 29 are curvature flags); along all 312 C2 sweeps A^* is nondecreasing and concave in the bound (0 failures), as A5 predicts.

TABLE III
I1 ERRORS $|\hat{\mu} - s|/|s|$ AT $\delta = 10\%$, CERTIFIED-BINDING CASES, BY f.

estimator	f	n	median	max	in $[L, H]$	med. $(H-L)/s$
big-M dual	0.25	7	2.6×10^{-1}	4.2×10^{0}	0	6.2×10^{-4}
big-M dual	0.5	8	1.4×10^{-1}	1.2×10^{0}	0	3.2×10^{-4}
big-M dual	0.75	8	1.9×10^{-2}	3.0×10^{-1}	0	3.2×10^{-4}
perspective	0.25	7	4.5×10^{-2}	8.2×10^{-1}	0	6.2×10^{-4}
perspective	0.5	8	8.2×10^{-3}	1.0×10^{-1}	0	3.2×10^{-4}
perspective	0.75	8	8.4×10^{-3}	5.2×10^{-2}	0	3.2×10^{-4}
fixed-integer	0.25	7	4.8×10^{-2}	7.8×10^{-1}	1	6.2×10^{-4}
fixed-integer	0.5	8	6.0×10^{-2}	8.2×10^{-1}	0	3.2×10^{-4}
fixed-integer	0.75	8	5.2×10^{-3}	2.9×10^{-2}	3	3.2×10^{-4}
convex baseline	all	24	3.7×10^{-4}	4.2×10^{-3}	–	–

B. Q2: Sensitivity under Integrality

Setup. I1 (Table II) perturbs 4 bounds by δ (sector and dollar at 2%, all styles at 0.10, the Mom bound at 0.025): 48 cases, 36 at $\delta = 10\%$. The estimators are the big-M relaxation dual, the perspective relaxation slope and the fixed-integer dual; the baseline (Sec. II-E) is the convex reference. The criterion is $|\hat{\mu} - s|/|s|$ over certified-binding cases, and whether $\hat{\mu} \in [L, H]$.

Certification. No interval is violated (0 of 144 points) and the formulations' polished values tie at all 144. At $\delta = 10\%$, 23 of 36 cases are certified binding (sector 9, momentum 6, dollar 5, all styles 3); in the others every estimator is zero (the perspective slope to within 1.4×10^{-11}). The median relative width $(H - L)/s$ is 3.6×10^{-4} (1.2×10^{-4}–0.020, $n = 23$): lying in $[L, H]$ means matching the secant to that width.

Estimator errors. Table III gives the errors at $\delta = 10\%$ ($n = 23$). The perspective slope has the smallest median error, 0.013 (max 0.82), then the fixed-integer dual, 0.023 (max 0.82), and the big-M dual, 0.073 (max 4.2). Only 4 estimates, all fixed-integer duals, lie in $[L, H]$. The baseline row counts the 24 cases where the baseline's own constraint binds ($e \ge 10^{-4}$): the 23 certified-binding cases and 1 more. Its dual is within a median 3.7×10^{-4} (max 4.2×10^{-3}) of its own secant, so integrality, not the step, causes the errors. The errors depend on δ: at 2017-12 ($n = 11$ binding at each step) the medians at $\delta = 5\%$ and 10% are 0.014 and 0.023 (fixed-integer), 0.052 and 0.016 (perspective), 0.24 and 0.073 (big-M).

When the fixed-integer dual is exact. Each case is classed by comparing, at $b - \delta$, b and $b + \delta$, the signed support of V's incumbent and the fixed-integer QP's active set: P1, both unchanged; P2, support unchanged, active set changed; P3,

support changed. Of the certified-binding cases, 3, 2 and 18 of 23 are P1, P2 and P3 at $\delta = 10\%$, and 4, 1 and 6 of 11 at $\delta = 5\%$. With fixed support and active set the QP solution is affine in b, so the value is quadratic and the central secant is the derivative at b. Over both steps, the fixed-integer error is at most 6.7×10^{-12} in P1 ($n = 7$), within the 10^{-6} predicted before the analysis ran, at most 2.9×10^{-4} in P2 ($n = 3$), and has median 0.056 in P3 ($n = 24$). The relaxation slopes fail differently: their error rises with the relaxation's gap at b (Spearman rank correlation 0.64 for the perspective, 0.50 for big-M, $n = 23$).

Value functions. I2 sweeps the sector bound, the I1 constraint with the largest median point slope over the 9 I1 instances at $\delta = 10\%$ (the design rule): 1.364×10^{-3} per unit of weight, against 1.316×10^{-3} per unit of loading for the Mom bound, a comparison across units; by median elasticity the Mom bound ranks first (0.027 against 0.020). The P-MI value lies below the baseline value and is not concave: 12 of 39 grid triples at 2017-12 and 9 of 39 at 2020-03 are certified nonconcave, $V(b^{(k-1)}) + V(b^{(k+1)}) > 2\bar{U}_{\text{eff}}(b^{(k)})$, and the support changes between 16 of 40 and 17 of 40 neighbouring grid points $b^{(k)}$. At $0.5b$, b and $1.5b$ ($\delta = 10\%$, $n = 6$; Fig. 3) the median errors are 0.32 (big-M), 0.078 (perspective) and 0.059 (fixed-integer), none in its interval (relative width at most 9.9×10^{-4}).

C. Q3: Formulations and Numerics

Setup. N1 (Table II) solves the 12 instances per tightness level (loose: styles 0.4, beta 0.2; base; tight: 0.025, 0.0125) with the three arms at MIPGap 10^{-4} and WorkLimit W^*. The criteria are Work and the root gap after Gurobi's root cut loop, (root bound − best known value)/|best known value|, the best known value being the largest polished V over the instance's arms and presolve settings. Default-scaling results describe how Gurobi treats the unscaled cones (solver behaviour), not the formulation.

Root gaps and effort. At default presolve and base tightness the rescaled perspective has the smallest median root gap, 3.3×10^{-4} against big-M's 4.4×10^{-4} (big-M's is smaller at tight; Fig. 4), and the least median Work at every level (base 0.66 against 1.0); the default-scaling perspective is far behind (0.015; Work 1.2). No solve at default presolve stopped at the WorkLimit. The plain relaxations order as A6 predicts on every instance (0 violations; median margin 0.057 of $|V|$), with median plain gaps 0.059 (big-M) and 1.8×10^{-3} (perspective) at base ($n = 12$). Presolve, not the formulation, closes big-M's gap: at $N = 250$ ($n = 12$ per arm), presolve off raises big-M's median root gap from 6.0×10^{-4} to 0.057 (plain: 0.060), and 6 of 12 solves stop at the WorkLimit, the others using 4.3–171 Work; the rescaled perspective moves only from 4.9×10^{-4} to 5.5×10^{-4} (median Work 0.766 to 2.28).

Two mechanisms. The default-scaling root bound is not the cone tolerance. N3 (Table II) repeats the root of both perspective arms at every N1 instance at default presolve, as root-only runs (*probes*); the default probe reproduces N1's root bound exactly at 36 of 36. Forcing outer approximation (MIQCPMethod 1; the default, −1, lets Gurobi choose) gives the same bound exactly at 29 of 36 (rescaled: 30), consistent with Gurobi choosing outer approximation at most instances; its logs are off, so the choice itself is not recorded. A QCP node relaxation (method 0) lowers the median default-scaling root gap from 0.015 to 5.7×10^{-4}, below the plain perspective relaxation's 1.9×10^{-3} over the same 36 instances, as root cuts allow (within $5.0 \times 10^{-3}|V|$ of it at 33), whereas the default's root bound is looser than the plain relaxation itself. FeasibilityTol 10^{-9} leaves it at 0.015: the ratio of the two gaps has median 0.98, within a factor 2 at 31 of 36 instances, 2 below, 3 above (max 4.4). The rescaled arm's gaps under the three methods stay within a factor 7.7 of each other. N3 was added after the other experiments had been analysed, to test on all 36 instances what the design pilot (S2d) had found on a subset of them. The incumbent mechanism is the cone tolerance (A9): at default presolve ($n = 36$ per arm) the default-scaling incumbents violate their cones within tolerance with an objective effect of median 4.6×10^{-4} and max 6.4×10^{-3} of $|V|$, at most 1.9×10^{-11} rescaled. The rescaled arm changes both the rows' scale (divided by $\max_i u_i$) and the cones' scale, so N1 does not separate the two. Every arm's polished value lies within every other arm's $\bar{U}_{\text{eff}}$ (0 failures in 144 pairs).

Objective scaling. N2(a) repeats N1's base instances at $N = 250$ at $s_{\text{obj}} = 1$ and 10^4. At $s_{\text{obj}} = 1$ big-M stops at the WorkLimit in 2 of 4 instances (the others use 2.9–4.7 Work), and the rescaled perspective at 2017-12 with $K = 67$ reported a closed gap while its $\bar{U}_{\text{eff}}$ lay $9.1 \times 10^{-5}|V|$ below big-M's polished value, a false certificate. At 10^4 no incumbent exceeds another arm's bound (0, against 1), and the default-scaling arm's median certified gap is 1.2×10^{-3}, above MIPGap 10^{-4}: its cone effect reaches 4.8×10^{-3} of $|V|$, while the rescaled arm's median gap is 9.0×10^{-5}. In N2(b), P-RB at 2008-10, 2013-06, 2017-12, 2020-03, 2025-04, scaling the objective by 10^3 cuts Gurobi's barrier errors from 3.1×10^{-5} to 7.2×10^{-7} in A^* and from 0.057 to 9.5×10^{-3} in the binding duals, while risk-row normalization (the risk cone divided by σ_m, to a unit right-hand side) leaves its A^* identical (10 of 10).

V. Discussion

TODO(Isaiah), V, 0.8 page; the coverage list of modeling choices, limitations, trade-offs, unexpected behavior and extensions (claim → evidence):

- Modeling choice, fixed risk budget against fixed λ: costs at equal risk; mean-variance recovered at λ^* (A1) → 3.8×10^{-8}.
- Modeling choice, λ^* fixed within (date, N) for P-MI: exact only without integers; the MIQP's volatility falls below the budget (I1 median 2.87%, all 48 of 48 base points below; N1 4 of 36 above), so the Q2 slopes are of the penalized objective → Sec. II, the baseline at 39 of 48 I1 base points on the budget.
- Modeling choices: single period with $w_0 = 0$; the six-factor, 252-day risk model with the residual floor; the momentum construction and winsorization; the base

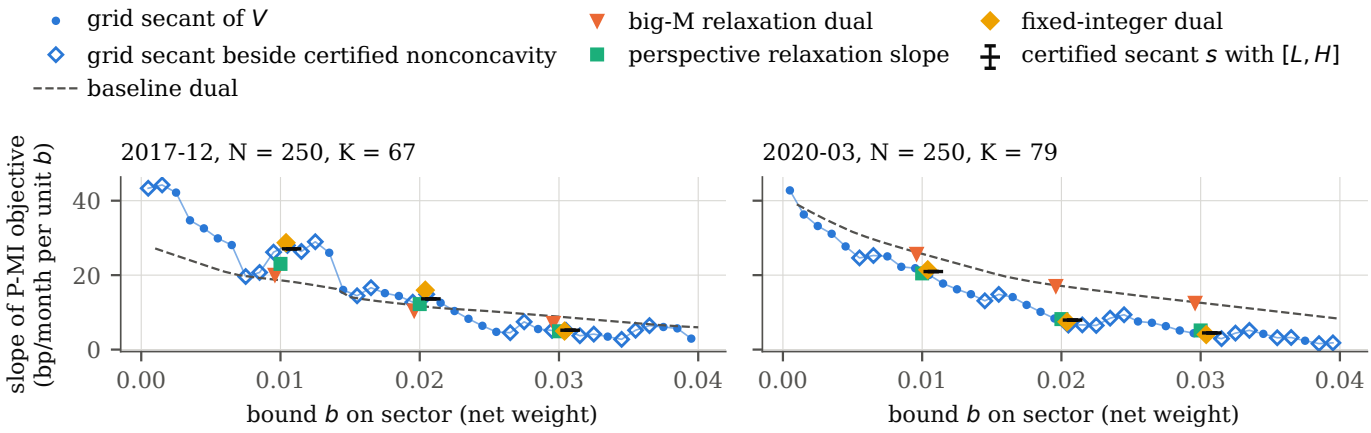


Fig. 3. I2 slopes on the sector bound (250 names; left 2017-12, right 2020-03), P-MI objective per unit of bound: neighbour secants of V (hollow: certified nonconcave), the baseline's dual, and the estimators beside s and $[L, H]$ at $0.5b$, b, $1.5b$ (offset).

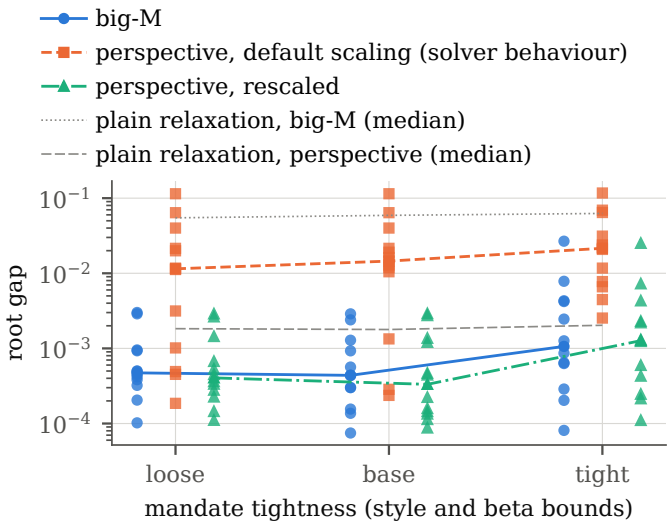


Fig. 4. N1 root gap against tightness, default presolve; $n = 12$ per cell; gray: median plain gaps.

mandate values; the 6% cap and ℓ in P-MI; costs as assumptions → Table I.

- Limitations: the disclosures below.
- Trade-off, alpha against exposure: the cost curves → Fig. 1, 0.744, 0.961.
- Trade-off, formulation strength against per-node cost: tighter at the root, but cone nodes → 0.057; median nodes at base 36.5 (big-M), 6.5 (rescaled), 270.5 (default).
- Trade-off, tightness against effort → Fig. 4, 1.1, 0.77.
- Trade-off, tolerance against runtime: MIPGap 10^{-6} (I1) against 10^{-4} (N1); at $s_{\text{obj}} = 1$ big-M stops at the WorkLimit in 2 of 4 (others 2.9–4.7 Work) → Sec. IV-C.
- What each estimator measures (A8), and why the errors depend on δ → 0.052 against 0.016 at 2017-12; at $\delta = 5\%$ the P1 maxima are 3.6 (perspective) and 5.4 (big-M), $n = 4$.
- Two failure modes: a support change within the step for the fixed-integer dual (P3 0.056; P1 6.7×10^{-12}), and relaxation looseness for the relaxation slopes (Spearman 0.64, 0.50; perspective per class 0.011, 0.014) → Sec. IV-B.
- The I2 near-tie: median secant 1.364×10^{-3} (sector, per unit of weight) against 1.316×10^{-3} (momentum, per unit of loading), a comparison across units; by elasticity the Mom bound ranks first (0.027 against 0.020); under the perspective slope the order flips, 1.222×10^{-3} against 1.306×10^{-3} → Sec. IV-B.
- Unexpected: presolve does big-M's root work → 6.0×10^{-4} against 0.057; 6 of 12 stopped at the WorkLimit with presolve off.
- Unexpected: the default-scaling root bound is consistent with outer approximation stopping early, not the cone tolerance → 29 of 36, 5.7×10^{-4}, 0.98.
- Unexpected: the false certificate at $s_{\text{obj}} = 1$ → 9.1×10^{-5}.
- Unexpected: the risk-row normalization is absorbed by presolve (supported, not proven) → 10 of 10, 1.6×10^{-5}.
- Unexpected: the hedged legs (in 7 of 45 perspective relaxations the relaxation splits a name into a long and a short leg to cheapen its cone; 11 name–relaxation pairs, 4 distinct names, the prediction holding at 11); per-leg cones as an extension → 11, 4, 11 (not yet in Results).
- Slope accuracy holds on every I1 case, not only the gated ones → 36 of 36.
- The earlier iteration's lessons in one or two sentences, without numbers: the fixed-risk formulation and the scaling policy.
- Extensions: platform coordination by exposure prices (decomposition); multi-period rebalancing (MPC); robust alpha (SOCP); a Lagrangian bound for the cardinality row.

VI. Conclusion

TODO(Isaiah), VI, 0.2 page (claim → evidence):

- Q1: neutrality is cheap for a weakly aligned alpha; the liquidity cap, sector bounds and trading costs carry the cost → 0.16, 0.16, 0.964.
- Q2: the fixed-integer dual is exact on a stable support and the best estimator depends on the step → 6.7×10^{-12}, 0.056, Table III.
- Q3: two mechanisms, outer approximation at the root (inferred) and the cone tolerance at the incumbent; the perspective formulation helps only in the rescaled formulation (rows and cones together), and objective scaling is still needed → 3.3×10^{-4}, 0.015, 29, 9.1×10^{-5}.
- Extensions → Sec. V.

Disclosures and AI-Use Acknowledgment

TODO(Isaiah), disclosures, the eight required items (claim → evidence):

- Universe: current S&P 500 constituents (snapshot 2026-10-16), so survivorship and look-ahead in the universe; GICS sectors are not point-in-time → results/data_report.md.
- Yahoo adjusted data; the JCI and DHR errors masked; a day is valid only with a return and dollar volume of at least $1M, a stock eligible with 80% valid days in its window; 1 multi-session return enters the regressions as a one-day return → results/data_report.md.
- One-pass winsorization leaves $|z|$ above 3: max 5.47 (2024-04), at 240 of 240 month-ends, 1936 of 107592 name-dates → results/disclosures/.
- Single period, $w_0 = 0$; borrow and trading costs are assumptions → Table I.
- No backtest and no performance claims; the alpha is a simple input signal.
- The parameter-setting rules of the design pilot (s_{obj}, W^*, the K grids, the I1 constraints) → Sec. III, experiments/configs/integer.yaml.
- The earlier iteration, without numbers.

TODO(Isaiah), AI-use acknowledgment, from docs/AI_USE.md's v2 log (claim → evidence):

- Claude (chat) designed v2 and extracted the repository; proposed most design decisions, check text, design edits and gate notes; sketched derivations; drafted the proposal text. Isaiah chose among the proposals, edited and signed every gate note, and ran the commands.
- Claude Code, with the result-checker and python-reviewer subagents, wrote and ran the code, tests and checks, built the figures and tables, and drafted Sections II–IV of this report (Mode A); Isaiah revises them and checks the derivations by hand.

References

[1] R. Grinold and R. Kahn, *Active Portfolio Management*, 2nd ed., 2000.
[2] R. Clarke, H. de Silva, and S. Thorley, "Portfolio constraints and the fundamental law of active management," *Financial Analysts Journal*, vol. 58, no. 5, pp. 48–66, 2002.
[3] R. Stubbs and D. Vandenbussche, "Constraint attribution," *Journal of Portfolio Management*, vol. 36, no. 4, pp. 48–59, 2010.
[4] J.-H. Lee and D. Stefek, "Do risk factors eat alphas?" *Journal of Portfolio Management*, vol. 34, no. 4, pp. 12–24, 2008.
[5] S. Ceria, A. Saxena, and R. Stubbs, "Factor alignment problems and quantitative portfolio management," *Journal of Portfolio Management*, vol. 38, no. 2, pp. 29–43, 2012.
[6] B. Jacobs, K. Levy, and H. Markowitz, "Portfolio optimization with factors, scenarios, and realistic short positions," *Operations Research*, vol. 53, no. 4, 2005.
[7] A. Frangioni and C. Gentile, "Perspective cuts for a class of convex 0–1 mixed integer programs," *Mathematical Programming*, vol. 106, no. 2, pp. 225–236, 2006.
[8] O. Günlük and J. Linderoth, "Perspective reformulations of mixed integer nonlinear programs with indicator variables," *Mathematical Programming*, vol. 124, pp. 183–205, 2010.
[9] A. Geoffrion and R. Nauss, "Parametric and postoptimality analysis in integer linear programming," *Management Science*, vol. 23, no. 5, pp. 453–466, 1977.
[10] L. Wolsey, "Integer programming duality: price functions and sensitivity analysis," *Mathematical Programming*, vol. 20, pp. 173–195, 1981.
[11] D. Bertsimas and R. Cory-Wright, "A scalable algorithm for sparse portfolio selection," *INFORMS Journal on Computing*, vol. 34, no. 3, 2022.
[12] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge University Press, 2004.
[13] L. Wolsey, *Integer Programming*, 2nd ed. Wiley, 2020.
[14] E. Fama and K. French, "A five-factor asset pricing model," *Journal of Financial Economics*, vol. 116, no. 1, 2015.
[15] K. R. French, "Data library."
[16] "yfinance."
[17] P. Goulart and Y. Chen, "Clarabel," 2024, arXiv:2405.12762.
[18] Gurobi Optimization, *Gurobi Optimizer Reference Manual*, 2026.
[19] S. Diamond and S. Boyd, "CVXPY," *Journal of Machine Learning Research*, vol. 17, no. 83, 2016.